

Deep Q-Learning with Rubik's Cubes

Thomas Kwashnak

May 7, 2023

Contents

1	Abstract	2
2	Background	3
2.1	Neural Networks	3
2.2	Deep Q-Learning	3
3	Project Structure	5
3.1	Rust	5
3.2	SQL	5
4	Analysis	6
5	Conclusions	7
A	SQL Code	8
A.1	Tables	8
A.1.1	Model	8
A.1.2	Epoch	8
A.1.3	Evaluation	8
A.1.4	EvaluationMove	9
A.1.5	Network	9
A.1.6	Bias	9
A.1.7	Weight	9
A.2	Users	10
A.2.1	Model	10
A.2.2	Reports	10
A.3	Views	11
A.3.1	EvaluationData	11
A.4	Procedures	11
A.4.1	get_current_epoch	11
B	Report Graph Source Code	12
C	GitHub Repository	13

1 Abstract

2 Background

2.1 Neural Networks

Neural Networks are one of the core pieces of machine learning. Neural Networks are complex formulas that you are able to train by providing inputs and expected outputs, and the model will try to best identify patterns in the data. Neural Networks are often used to classify images, text, and many other applications.

In short, a Neural Network is a function approximator. It achieves this through use of gradient descent with its many variables. Neural Networks may have upwards to 100,000s of variables, consisting of weights and biases. To train these weights and biases, you need to provide pre-labeled data, or data that already has a predetermined answer.

First, the input data is fed into the network. The network will pass the data through the layers, and eventually output what it thinks the answer is. You then compare the result to the expected value, and square the result. The result will be your “loss” for that observation. Loss tracks how badly the network did in this scenario.

With loss, and the loss function, we can calculate the gradient for the entire function. The gradient consists of the partial derivatives for every single weight and bias in the network. Recall that these weights and biases are what determines the output. By finding the partial derivative of each variable as it relates to loss, we can take use the negative gradient to nudge each variable so that the loss decreases.

By performing this sequence multiple times, the network will start figuring out patterns in the data. Eventually, we can throw “unlabeled data” and see what it thinks about it. Additionally, there are some configurations that we can make with this process. We call these parameters “hyperparameters”. Hyperparameters are parameters that appear to be arbitrary. In truth, many times we can’t fully understand the effects of one hyperparameter on the overall process. As such, very little specific guidance is provided regarding how to set these variables. For a basic neural network, for example, hyperparameters could be the number of layers it has, how many nodes the network has in each layer, or the rate at which it learns¹.

2.2 Deep Q-Learning

Deep Q-Learning is a form of Reinforcement Learning that utilizes a Q-Function to train the model against. The goal of Deep Q-Learning is to explore an unknown environment with a clear goal, learning from prior experience to maximize reward.

Deep Q-Learning requires the user of an environment. In short, an environment contains observations, actions the agent can take, and some form of a reward system that drives the user towards the end goal. For this project, the environment is simply a Rubik’s cube². The observations consist of an array of values representing each tile face of the cube³, the actions

¹The value we multiple the negative gradient by before applying it to all of the variables

²Originally the environment was a standard 3x3 Rubik’s Cube, however to simplify things I’ve recently switched over to a 2x2 Rubik’s Cube, with intentions of returning to a 3x3 once I create a model that solves a 2x2 Rubik’s Cube

³On a 3x3 Rubik’s Cube, this would be each of the 9 smaller squares on each side of the cube

are each of the possible moves that can be performed on the cube, and the reward is the number of correct squares on each face of the cube⁴.

In Deep Q-Learning, the logic that chooses the answer is called a Q-Function. The Q-Function consumes the current state of the cube, and returns the predicted reward it thinks each move will give it. For example, if the cube is one move away from being solved, the reward for the move that solves the cube will be higher than all other moves.

In standard Q-Learning practice, the Q-Function is a table that records all states that the agent has seen, what actions they took, and what the resulting reward actually was. However, this is unfeasible for environments such as the Rubik's Cube. It is estimated that the Rubik's Cube has 43 Quintillion possible states. To make sense of that number, if you scramble the cube randomly, it is statistically impossible to reach a state that has already been reached before. Thus, a table isn't effective.

That's where Neural Networks come in! Neural Networks are great at approximating functions, like Q-Functions! In our implementation of Q-Learning we will use 2 Neural Networks. One is called the Q-Network, and the other will be called the Target network.

But why 2 networks? Aren't we only training one? Yes. We are only training one, but one of the problems with reinforcement learning is that there is no "answer" that we can train the network on. Thus, we need to make that answer. The goal with training is to make the Q-Network approximate the predicted reward of taking an action at State A to be equal to the reward of State B plus the Target Network's approximation for the best move from State B.

That... probably didn't make sense. It's a very complicated subject, and it took me a long time to figure out. Imagine you're playing chess, and you're looking 2 moves ahead. You then associate the first move with the reward you get after the following move. Similarly, we are training the Q-Network to associate an action with the actual reward plus the predicted future reward of that state⁵.

In order to keep the target network up to date, every so often we take the target network and update it with the values of the Q-Network. In a perfect world, where the Q-Network perfectly predicts the future target network's result, every time we update the target network, the Q-Network will be looking one more step ahead.

There is a lot I left out in this background section, and a lot that was probably confusing. However, much of the analysis relies on some basic knowledge included here, but shouldn't be too difficult to understand once you figure out the basics.

⁴The reward function has been a tough point to figure out with this project. For now this method will work, but there are surely better reward systems out there

⁵And if there are people who are still confused at this point, I don't know how else to word it.

3 Project Structure

Currently, the project is split up into individual components. The following sections discuss the uses for each component and how they fit into the bigger picture. Originally, this project was just a large project written in python, but I have since split it up during a complete re-write that started at the beginning of the Spring 2023 Semester.

3.1 Rust

In prior iterations of the project, I was never quite satisfied with how efficient the emulation of the Rubik's Cube was in python. Many times it ended up being a bottle-neck for the entire system. Thus, I decided that I would write the environment in Rust.

This project uses the PyO3 library for Rust, which provides the ability to compile Rust into a Python module, allowing me to simply import and use it directly in python. The benefit of this is that all of the logic performed with the cube is done in a memory-safe, type-safe, efficient language that runs at the speed of C.

Furthermore, I decided to also implement the replay database in Rust, giving me even more optimizations when it comes to storing and fetching samples to train on.

3.2 SQL

I made some mistakes in the original implementation. Originally, I was using large JSON files to store all of the collected data for analysis. This had a few negative effects. Firstly, the data was unoptimized. Secondly, it required me to load the entire file into ram every time I wanted to add another value⁶. Lastly, there were some additional constraints with the storage of numbers within data file.

Thus, to solve these problems I decided to use SQL to store everything. I set up and configured a Microsoft SQL Server running on a Docker Container. This solved many of the problems above, and also had the additional benefit of being accessible by any of my devices⁷.

⁶This quickly built up over time, especially when I started leaving the program running overnight and throughout all of my classes

⁷This is achieved through the use of Tailscale, which creates a private virtual network between all of your registered devices

4 Analysis

5 Conclusions

A SQL Code

A.1 Tables

A.1.1 Model

This table is used to differentiate between different models, as each model will have different configurations, sizes, and other features that need to be kept separate. This table consists of both historical records parsed from JSON and new models that are used in current analysis.

```
CREATE TABLE Model
(
    ModelId    int identity primary key,
    ModelName  VARCHAR(100),
    GitHash    VARCHAR(40),
    CubeType   VARCHAR(50)
)
```

A.1.2 Epoch

The Epoch table stores all of the epoch data collected during training

```
CREATE TABLE Epoch
(
    EpochId int identity primary key,
    ModelId int not null foreign key references Model (ModelId),
    Epoch   int not null,
    Loss    float(53),
    Reward  float(53)
)
```

A.1.3 Evaluation

The Evaluation table stores related data for evaluation runs by the model. The primary function of the Evaluation table is to keep a record of when the model starts making varied moves (instead of repeating the same move over and over).

```
CREATE TABLE Evaluation
(
    EvaluationId int identity primary key,
    ModelId      int not null foreign key references Model (ModelId),
    Epoch        int not null,
    Solved       bit,
    MoveCount     int,
    Seed         bigint,
)
```

A.1.4 EvaluationMove

The Evaluation Move table contains each individual move made during evaluations. Each move is tied to an evaluation, a move index, and stores the move and its reward.

```
CREATE TABLE EvaluationMove
(
    EvaluationId INT NOT NULL FOREIGN KEY REFERENCES Evaluation
    → (EvaluationId),
    MoveIndex    INT NOT NULL,
    MoveName     VARCHAR(10),
    Reward       FLOAT(53),
    PRIMARY KEY (EvaluationId, MoveIndex)
)
```

A.1.5 Network

Acts as an identification item for each unique network stored in the Bias and Weight tables.

```
CREATE TABLE Network
(
    NetworkId INT IDENTITY PRIMARY KEY,
    ModelId   INT FOREIGN KEY REFERENCES Model (ModelId),
    Epoch     INT,
    IsTarget  BIT NOT NULL DEFAULT 0
)
```

A.1.6 Bias

Stores each of the biases within a network.

```
CREATE TABLE Bias
(
    NetworkId INT NOT NULL FOREIGN KEY REFERENCES Network (NetworkId),
    Layer     INT NOT NULL,
    X         INT NOT NULL,
    Bias      FLOAT(53),
    PRIMARY KEY (NetworkId, Layer, X)
)
```

A.1.7 Weight

Stores each of the weights within a network.

```
CREATE TABLE Weight
(
    NetworkId INT NOT NULL FOREIGN KEY REFERENCES Network (NetworkId),
```

```

    Layer      INT NOT NULL,
    X           INT NOT NULL,
    Y           INT NOT NULL,
    Weight      FLOAT(53),
    PRIMARY KEY (NetworkId, Layer, X, Y)
)

```

A.2 Users

A.2.1 Model

The user used when running and training the model itself. The Model user must be able to insert and update data, as well as remove old networks to preserve space.

```
CREATE LOGIN Agent WITH PASSWORD = 'MlCubeAgentPass1234';
```

```
CREATE USER Agent FOR LOGIN Agent;
```

```
GRANT CONNECT TO Agent;
```

```
GRANT INSERT ON Model TO Agent;
```

```
GRANT SELECT ON Model TO Agent;
```

```
GRANT SELECT ON Network TO Agent;
```

```
GRANT INSERT ON Network TO Agent;
```

```
GRANT SELECT ON Weight TO Agent;
```

```
GRANT INSERT ON Weight TO Agent;
```

```
GRANT SELECT ON Bias TO Agent;
```

```
GRANT INSERT ON Bias TO Agent;
```

```
GRANT INSERT ON Evaluation TO Agent;
```

```
GRANT INSERT ON EvaluationMove TO Agent;
```

```
GRANT INSERT ON Epoch TO Agent;
```

```
GRANT EXECUTE ON get_current_epoch TO Agent;
```

```
GRANT EXECUTE ON delete_network TO Agent;
```

A.2.2 Reports

The user used when building reports and graphs. This user only needs to be able to select and fetch data from various tables.

```
CREATE LOGIN Reports WITH PASSWORD = 'mlcube-reporting123'
```

```
CREATE USER Reports FOR LOGIN Reports;
GRANT CONNECT TO Reports;
```

```
GRANT SELECT ON Model TO Reports
GRANT SELECT ON Epoch TO Reports
GRANT SELECT ON Evaluation TO Reports
GRANT SELECT ON EvaluationMove TO Reports
GRANT SELECT ON GroupedEpoch TO Reports
GRANT SELECT ON EvaluationData TO Reports
```

A.3 Views

A.3.1 EvaluationData

A collected table from both the Evaluation and EvaluationMove tables. Used for fetching data for reports and analysis.

```
CREATE VIEW EvaluationData AS
SELECT ModelId, Epoch, Evaluation.EvaluationId Id, Solved, Seed,
    ↪ MoveIndex, MoveName, Reward
FROM Evaluation
    LEFT JOIN EvaluationMove ON Evaluation.EvaluationId =
    ↪ EvaluationMove.EvaluationId
```

A.4 Procedures

A.4.1 get_current_epoch

Returns the current epoch for a specified model. Calculates by grabbing the last epoch recorded for that model.

```
create procedure get_current_epoch(@ModelId int)
AS
begin
    SELECT E.Epoch
    FROM Epoch E
    WHERE E.ModelId = @ModelId
        AND Epoch = (SELECT MAX(Epoch)
                     FROM Epoch
                     WHERE Epoch.ModelId = @ModelId)
end
```

B Report Graph Source Code

The following is the source code used to create the graphs

```
library(tidyverse)
library(ggplot2)
source("src/database.R")
source("src/functions.R")

file_name <- function(name) {
  return(paste("../SER-300-Report/assets/", name))
}

save_gg <- function(name, height = 4) {
  ggsave(file_name(name), width = 6.5, height = height, units = "in")
}
```

C GitHub Repository

All source files, including the ones found in prior appendixes, can be found on the GitHub Repository. <https://www.github.com/LittleTealeaf/mlcube>